

A Statistical Machine Translation System for Spanish-to-English Translation

Albert Chu, Brad Huang, and Nicholas P. Moores, Stanford University

Dan Jurafsky (for CS 124)

Word Count: 2388

(Dated: 2 March 2015)

I. COMMENTS ON *F*, SPANISH

Please see the appendix for examples referred below.

To enumerate how we were able to capitalize on or were challenged by building this machine translation system, it is necessary to first discuss the relevant differences between English and Spanish. Key differences:

1. Spanish is much more heavily inflected than English, with different verb conjugations for *-ir* and *-er* and *-ar* final verbs and different conjugations for mood and tense. English only inflects verbs for number in the third-person present tense, and has largely abandoned marking verbs for number in the past tense.
2. Within noun phrases, Spanish adjectives tend to appear following the nouns they modify instead of before them as in English (though adjective placement in Spanish can be a cue to meaning)
3. Spanish word ordering is much more flexible than is English's. Whereas English rarely deviates from SVO (Subject-Verb-Object) word ordering, Spanish speakers do so whenever they desire to place particular emphasis on a word (and will then place that word toward the end of the sentence).
4. Spanish uses grammatical gender, whereas English abandoned such a system hundreds of years ago.
5. Spanish heavily relies on the use of the subjunctive mood when describing the past, whereas the use of the subjunctive mood in English has become less commonplace.

6. Spanish nouns and adjectives must also agree by both gender and number, as must the particles used at the beginning of noun phrases (for example, the definite articles *el*, *la*, *los*, and *las*).
7. English indefinite articles are marked for whether their noun begins (phonetically) with a vowel, whereas Spanish does not.
8. Spanish will often use a (non-progressive) present tense form to indicate that something is currently happening, whereas English will use primarily the progressive present tense form to indicate this; a non-progressive translation would seem wrong to a native English speaker.

Many of the difficulties that arise from translating Spanish into English rest on the morphosyntactic differences between the two languages. The lack of grammatical gender in English poses the problem that many of the words in a Spanish sentence that agree based on the gender of the noun at the head of the noun phrase must instead map to the same English articles and adjectives (but that the same does not hold for number). For example, it is therefore a challenge to train a system to recognize that *traducido* and *traducida* map to the same English word, *translated*. Fortunately, since it is Spanish that is more inflected, it is easier to translate from Spanish to English since one can generalize across or collapse inflections, whereas an English-to-Spanish translation would require the system to correctly output sentences whose words are matched for Spanish gender, number, tense, and mood.

II. IBM IMPROVEMENT STRATEGIES

In the course of this project we applied several separate strategies attempting to improve our statistical machine translation system.

1. `LanguageModel.py`

We first implemented a trigram language model with stupid backoff on English. For each Spanish word f to be translated, we select the top two words with the highest $t(f|e)$ probabilities. However, for a Spanish sentence of length m , we would have 2^m combinations of possible word-for-word translations. To limit the search space, the variable α controls how many second-highest probability words can appear in one sentence, and is usually set at $\alpha = 5$ for development.

2. `PhraseTable.py`

We then attempted to further improve the model by implementing a phrase translation. During training, the IBM Model 1 is run twice to construct probability tables from English to Spanish and from Spanish to English. Word pairs of reciprocal best word-to-word translation probability are then selected, and some other word pairs are selected heuristically. Smaller phrases are then combined into larger phrases and a dictionary mapping Spanish phrases to English phrases is thus constructed. The probability of mapping Spanish phrases to English phrases is computed based on the counts of these phrases. Possible English sentences constructed are scored by the trigram language model and the probabilities of phrase mappings.

However, we discovered that the phrase table does not increase either the BLEU-1 or BLEU-2 scores. One common type of errors phrase table makes is that words that appear less frequently in both languages may not be combined into phrases, and thus may be eliminated from the translation of phrase table, thus decreasing the BLEU scores. (Output of this script is named `translations.pt` in the output folder for reference.) Using a much larger training set could help remedy this, so that sufficient phrase mappings can be found.

3. `nltpos.py`

Next, we used post-processing to improve the translation. One common word order difference between English and Spanish is that Spanish sentences usually have nouns

preceding their modifiers while English is the opposite. As a result, we use NLTK’s POS tagger to identify noun-adjective pairs and reverse the order. This increases the BLEU-2 score while not affecting the BLEU-1 score as expected since BLEU-1 is based on correct unigrams and BLEU-2 is based on correct bigrams.

4. *Postprocessing*

We introduced several other post-processing strategies. We observed that because of the unweighted alignments considered in IBM Model 1, some words are translated into stop words, so we check for and eliminate consecutive stop words (which should not appear in English) after translation.

In addition, since English requires phonological agreement between articles and nouns (i.e. ‘a dog’ vs ‘an elephant’), we also check for this agreement and change the articles to ‘a’ or ‘an’ accordingly.

For those Spanish proper nouns not captured in the IBM model, they are taken as is in translation. This improves BLEU1 score by about 1.5 points.

Finally we attempted to correct capitalization of nations and months and the first letter in a sentence after translation. However, BLEU scores are not affected.

For reference, the following is a summary of BLEU scores:

BLEU scores on Dev Set

	Baseline IBM Model 1	Language Model	Phrase Translation
BLEU-1	47.68	51.74	43.36
BLEU-2	9.18	12.43	7.47

BLEU scores on Test Set

	Baseline IBM Model 1	Language Model	Phrase Translation
BLEU-1	49.23	51.36	45.35
BLEU-2	9.59	11.59	8.59

III. ERROR ANALYSIS

Through the course of the project, we implemented two improvements to the baseline model: A trigram language model with stupid backoff on English, and implemented phrase translation. We report the characteristics of our errors as they occurred in the baseline case, in the case where the system utilized the language model, and in the case where the system utilized phrase translation.

A. Version-specific Error Analysis

Baseline

Both infrequent Spanish words and infrequent inflected forms are often translated as *fawn*. In addition, proper nouns are encountered so infrequently that they are often replaced with *fawn* in the translation or are omitted altogether. One improvement is to use a gazeteer as additional input to the system to help correct errors in handling proper nouns.

This behavior happens since all alignments are weighted equally. One solution would be to limit the alignment to a few words before and after each word, so that in the sentence *yo voy a la tienda*. (translation: *I go to the store.*), *voy* could align to ± 1 precision (only aligning to *I* and *go* and *to*).

One of the biggest errors was that since there are so many inflected forms, some of them are not identified as the same word as the others. One method we could use in the future is to try stemming the source text.

Language Model

The errors here relate to improperly ordered translation of noun phrases. Since Spanish nouns usually precede the adjectives that modify them, the English translations given by this model mistakenly keep the original order of the Spanish words, failing to correct for the different ordering of noun phrases. Our system correctly translates the Spanish *de* to English *of*, but fails to take the adjective which follows *de* in the Spanish source sentence and place it directly before the English noun in the translated sentence which could be solved by identifying adjectives and placing them in front. In these instances, the output is semantically correct but syntactically uncommon. The English translated counterparts for Spanish stop words also often end up being erroneously removed as our model has trouble

correctly reinstating the stop words in the English translation which could be helped by limiting consecutive stop words or using NLTK treebank.

Phrase Translation

Longer sequences of words are correct when we use phrase translation, but the English translations are often missing stop words or entire dependent clauses. As a whole, our system with phrase translation maintains proper word order but omits words that belong in a correct translation. We could address these errors by encouraging the system not to eliminate infrequent words from the translation of the phrase table.

B. General Error Analysis

In general, most errors occur due to the following nuances of the difference between English and Spanish:

1. **Length:** Spanish sentences are generally longer than English sentences, so it is beneficial to reduce the length of the sentence when translating by restricting the length of the target sentence and eliminating words in the source sentence based on rankings like the Moldovan et al. criteria.
2. **Alignment:** All possible alignments of given sentence pairs are treated equally. This means that the probability of aligning repeating words to words in another language is abnormally high. As a result, translations from the baseline contains unrelated words that occur only because they are repeated several times in many sentences of the corpus. Some ways to improve would be: (1) make sure that a word does not align to words far from its corresponding index in the other sentence, (2) use stop word filtering to filter out abundant words like *el*, *la*, *the*, *a*, et cetera, and to assign hard rules to translate these stop words, (3) use a part-of-speech (POS) tagger to tag the training documents, and only consider alignments that align words of the same POS. Methods (1) and (3) are difficult because they increase the runtime a lot due to increased computational complexity (we have to actually look into each alignment), while method (2) would create some issues because there may not be 1-some correspondences between all Spanish and English stop words.
3. **Word Order:** Spanish and English have in general the same word orders; however,

Spanish word order is more flexible than English word order. To properly account for this, we would need a syntactic language model. See the second Google translate output as an example.

IV. COMPARATIVE ANALYSIS WITH GOOGLE TRANSLATE

Please see the appendix for a full enumeration of the source sentence, our output, and the gold translations.

1. **Google Translation:** This attorney for 48 years specializing in human rights and president of the Egyptian Organization for Human Rights (EOHR) defended the Muslim brothers when they were held in prison or had to deal with a trial during the dictatorship of Mubarak.

Analysis:

- Our translation translates the age of the lawyer ('de 48 años') correctly in a roundabout way though Google does not.
 - Spanish omits subjects whenever possible and uses inflections to identify subjects mentioned before. Our translation did not account for verb inflections, so we were not able to identify that the subject for '... se encontraban...' is 'hermanos musulmanes' as mentioned right before that verb.
 - From this translation we can observe that our model is poor at translating verbs, a result of the various inflected forms possible in Spanish. Google probably encodes Spanish inflections in their translation model, so that they can identify subjects and translate words more successfully, and as a result Google's translation is much more understandable.
2. **Google Translation:** "But oil is not the only reason why it is going so well" Says Anna, our waitress, while passing trays rakfisk and, with her long blonde hair and striking blue eyes, represents the image of the Norwegian welfare.

Analysis:

- In the quote section, the latter part is actually 'why we are doing so well', and neither translations are perfect in this case. Google changes 'nos' into 'it' because of the following 'va' (third person of 'ir'), while our translation just did word-to-word translation to get 'we is'.

- Although the Spanish aligns well to the actual translation in terms of word order, both Google and our translation translated the word ‘representa’ while that word is not in the actual translation.
- There are several deletions in our translation that should not occur - ‘our (waitress)’ and ‘long (blonde hair)’. This occurs to words with low frequency in the corpus. Because their appearances are so infrequent, they are mapped to NULL during IBM Model 1’s training process, and thus eliminated from final translation. As a result, Google’s translation ends up being more understandable.

3. **Google Translation:** Men and women mention his career, stress and cost of property and education as reasons that prevent them from having children.

Analysis:

- ‘su’ in Spanish gets translated into ‘his’ by Google (true in a word-to-word scenario), but ‘their’ by our model (better here, since it refers to both ‘Mena and women’), though it should actually be ‘sus’ here since the phrase refers to ‘their careers’.
- Although we tried to have a good statistical strategy to eliminate stop words, our model does place multiple ‘the’s in the final translation.
- Similarly, when the Spanish has ‘... el costo *de* las preiodades y *de* la educacin ...’, our model translated both ‘de’ in the sentence, while Google identified that the sentence does not require both.
- ‘motivos’ and ‘impiden’ are translated incorrectly, one into a stop word and another eliminated in our translation. As mentioned above, these inflected forms likely appeared too few times in the corpus, which led to incorrect word-to-word mappings. Google has the better translation for being more understandable.

4. **Google Translation:** The measures taken since 2009 ensured that the rate of registration of new voters fell in 2010 compared with 2006.

Analysis:

- Google translated ‘lograron’ into ‘ensured’ while our model translated it into ‘succeed’. ‘Lograr’ usually means ‘achieve, succeed’, and ‘lograron’ here actually means ‘lead to’ in this case.
- ‘inscripcin’ and ‘cayera’ are incorrectly translated likely due to the same reasons in the previous examples. Google’s translation is fairly more understandable.

5. **Google Translation:** The Prime Minister recently said the ODS (Civic Democratic Party) would not bother entrepreneurs with inspectors. So what is prohibited or permitted?

Analysis:

- A clear difference between our translation and Google’s translation is the negation. Google identified that there is negation in the sentence and translated accordingly, while our model did not.
- In the same vein, Google identified the inflected form of the verb and translated into ‘would not bother’, while our model does nothing about Spanish inflection, so our translation does not align to Spanish as well as Google.
- It is interesting to observe that in the last part of the sentence, where the real translation is ‘- so is it forbidden or allowed ?’, both Google and our model failed to identify that the subject of the question is ‘ODS’.

V. APPENDIX

A. Illustrating Examples of Language Differences

1.
 - Spanish: *el hombre corre* and *yo corro* and *los estudiantes corren*
 - English: *the man runs* and *I run* and *the students run*
2.
 - Spanish: *un pobre hombre* and *un hombre pobre*
 - English: *a poor (unfortunate) man* and *a poor (destitute) man*
3.
 - Spanish: *Masoud trabaja en el mercado.* and *Trabaja Masoud en el mercado.* and *En el mercado trabaja Masoud.*
 - English: *Masoud works in the market.*
4.
 - Spanish: *la abuela hermosa* and *el abuelo hermoso* and *los estudiantes hermosos*
 - English: *the beautiful grandmother* and *the beautiful grandfather* and *the beautiful students*
5.
 - Spanish: *Cuando eras estudiante, tu no trabajabas mucho* and *Cuando eras estudiante, no trabajaste mucho*
 - English: *When you were a student (for a long time), you didn't work much (habitually)* and *When you were a student (for a long time), you didn't work much (for a set time)*
6.
 - Spanish: *una abuela* and *una computadora*
 - English: *a question* and *an answer*
7.
 - Spanish: *una abuela* and *una computadora*
 - English: *a question* and *an answer*
8.
 - Spanish: *Ellos comen* but *Ellos son comiendo.**
 - English: *They eat** and *They are eating*

B. System Output as Compared to Google Translate Output Examples

1. **Spanish Source:** Este abogado de 48 años especializado en derechos humanos y presidente de la Organizacin Egipcia de Derechos Humanos (EOHR) defendi a los hermanos musulmanes cuando se encontraban retenidos en prisin o tuvieron que hacer frente a un juicio durante la dictadura de Mubarak .

Google Translation: This attorney for 48 years specializing in human rights and president of the Egyptian Organization for Human Rights (EOHR) defended the Muslim brothers when they were held in prison or had to deal with a trial during the dictatorship of Mubarak.

Our Translation: This lawyer of 48 years specialises in rights human and President of the organisation Egyptian of human rights (EOHR) for to the Muslim Muslims when who held in prison or had that to with to a trial for the dictatorship of Mubarak .

Actual Translation: The 48-year-old human rights lawyer and Chairman of the Egyptian Organisation for Human Rights (EOHR) defended the Muslim Brotherhood , when imprisoned or in court under Mubarak .

2. **Spanish Source:** " Pero el petróleo no es la única razón por la que nos va tan bien " , afirma Anna , nuestra camarera , mientras va pasando bandejas de rakfisk y que , con su larga melena rubia y sus sorprendentes ojos azules , representa la imagen del bienestar noruego .

Google Translation: " But oil is not the only reason why it is going so well " Says Anna, our waitress, while passing trays rakfisk and, with her long blonde hair and striking blue eyes, represents the image of the Norwegian welfare.

Our Translation: " but the oil is not the only reason the we is so well " , says Anna , our , while is happening explained of and that , with their long and their surprising blue eyes , represents the image the welfare Norwegian .

Gold Translation: " But oil is not the only reason we are doing so well , " says Anna our waitress , handing round trays of maturing rakfisk and , with her long blond hair and startlingly blue eyes , the image of Nordic well-being .

Actual Translation: " But oil is not the only reason we are doing so well ,

" says Anna our waitress , handing round trays of maturing rakfisk and , with her long blond hair and startlingly blue eyes , the image of Nordic well-being .

3. **Spanish Source:** Hombres y mujeres mencionan su carrera , el estrés y el costo de las propiedades y de la educación como los motivos que les impiden tener hijos .

Google Translation: Men and women mention his career, stress and cost of property and education as reasons that prevent them from having children.

Our Translation: Men and women mention their career , the stress and the costs of the properties and of the education as the for that them from have children .

Actual Translation: Men and women alike mention their careers , stress and the cost of property and education as the reasons preventing them from having children .

4. **Spanish Source:** Las medidas adoptadas desde 2009 lograron que la tasa de inscripción de nuevos electores cayera en 2010 en comparación con el año 2006 .

Google Translation: The measures taken since 2009 ensured that the rate of registration of new voters fell in 2010 compared with 2006.

Our Translation: The measures taken since 2009 succeed that rate of for of new electorate out in 2010 in compared with the year 2006 .

Gold Translation: However , the measures adopted since 2009 have led to a 17 % drop in the registration rate of new voters in 2010 compared to 2006 .

Actual Translation: However , the measures adopted since 2009 have led to a 17 % drop in the registration rate of new voters in 2010 compared to 2006 .

5. **Spanish Source:** El primer ministro dijo recientemente que la ODS (Partido Democrático Cívico) no molestaría a los empresarios con inspectores . Entonces , est prohibido o permitido ?

Google Translation: The Prime Minister recently said the ODS (Civic Democratic Party) would not bother entrepreneurs with inspectors. So what is prohibited or permitted?

Our Translation: The first Minister said recently that ODS (democratic civic party) not to the employers with inspectors . then , is ban or allowed ?

Actual Translation: The Prime Minister recently claimed that the ODS will not be burdening business owners with its checks - so is it forbidden or allowed ?